

# **Game of Life Report**

## **1. PCAM**

Partitioning - We decompose this problem into smaller subdomains.

To do that, my process involved decomposing the problem into smaller subdomains. We divide cell grid by columns

Communication - I am using 4 processors. Also MPI to exchange data about ghost cells. Each processor handles a partitioned area

Agglomeration - Each partitioned region is grouped into cells. I also use the gather MPI function to collect data from all 4 processes and put it together

Mapping - Make sure that tasks are equally distributed to processors.

In my 20 x 20, each of the 4 processors gets its fair share

## **2. My fortran code file is in gol.f90 and gol2.f90.**

To compile my file(s), write this command: `mpif90 gol(2).f90 -o gol(2).`

To run my file, write this command: `mpirun -np <Number of Processors> ./gol(2).`

To adjust the grid dimensions, modify `nrows` and `ncols`.

To adjust the number of steps, adjust the value of `N`

Also, you can uncomment the section of my code where I assign pseudo random boolean values to the grid.

The difference between the 2 files is that the `gol.f90` file is where I do question 3, and track the alive cells, whereas the `gol2.f90` file is where I do column-wise domain decomposition.

Below, I have snippets with comments explaining the purpose of each of those snippets:

```

! Set the rows
game(1, :) = game(nrows + 1, :)
game(nrows + 2, :) = game(2, :)

! Set the columns
game(:, 1) = game(:, ncols + 1)
game(:, ncols + 2) = game(:, 2)

! Set the corners
game(1, 1) = game(nrows + 1, ncols + 1)
game(1, ncols + 2) = game(ncols + 1, 2)
game(nrows + 2, 1) = game(2, ncols + 1)
game(nrows + 2, ncols + 2) = game(2, 2)

```

```

do i = 2, nrows + 1
  do j = 2, ncols + 1
    true_count = 0

    ! Checks all the cells around the current one
    do mini_i = i - 1, i + 1
      do mini_j = j - 1, j + 1
        if (mini_i /= i .or. mini_j /= j) then
          if (game(mini_i, mini_j)) then
            true_count = true_count + 1
          end if
        end if
      end do
    end do

    ! Game of Life Rules
    if (true_count .eq. 3) then
      next_game(i, j) = .TRUE.
    else if (true_count .eq. 2) then
      continue
    else
      next_game(i, j) = .FALSE.
    end if
  end do
end do

```

```

! Column-wise domain decomposition starts here
local_ncols = ncols / size

! Allocating memory for grids and communication buffers
! Each of the 4 processes will use its own grid
allocate(local_game(nrows + 2, local_ncols + 2))
allocate(local_next_game(nrows + 2, local_ncols + 2))
allocate(send_left(nrows + 2))
allocate(recv_left(nrows + 2))
allocate(send_right(nrows + 2))
allocate(recv_right(nrows + 2))
allocate(send_up(local_ncols + 2))
allocate(recv_up(local_ncols + 2))
allocate(send_down(local_ncols + 2))
allocate(recv_down(local_ncols + 2))

```

```

! If there is > 1 processors, exchange info about grid ghost cells
if (size > 1) then
  send_left = local_game(:, 2)
  send_right = local_game(:, local_ncols + 1)
  send_up = local_game(2, :)
  send_down = local_game(nrows + 1, :)

  call MPI_SENDRECV(send_left, nrows + 2, MPI_LOGICAL, left, 1, &
    recv_right, nrows + 2, MPI_LOGICAL, right, 1, &
    MPI_COMM_WORLD, MPI_STATUS_IGNORE, ierr)

  call MPI_SENDRECV(send_right, nrows + 2, MPI_LOGICAL, right, 2, &
    recv_left, nrows + 2, MPI_LOGICAL, left, 2, &
    MPI_COMM_WORLD, MPI_STATUS_IGNORE, ierr)

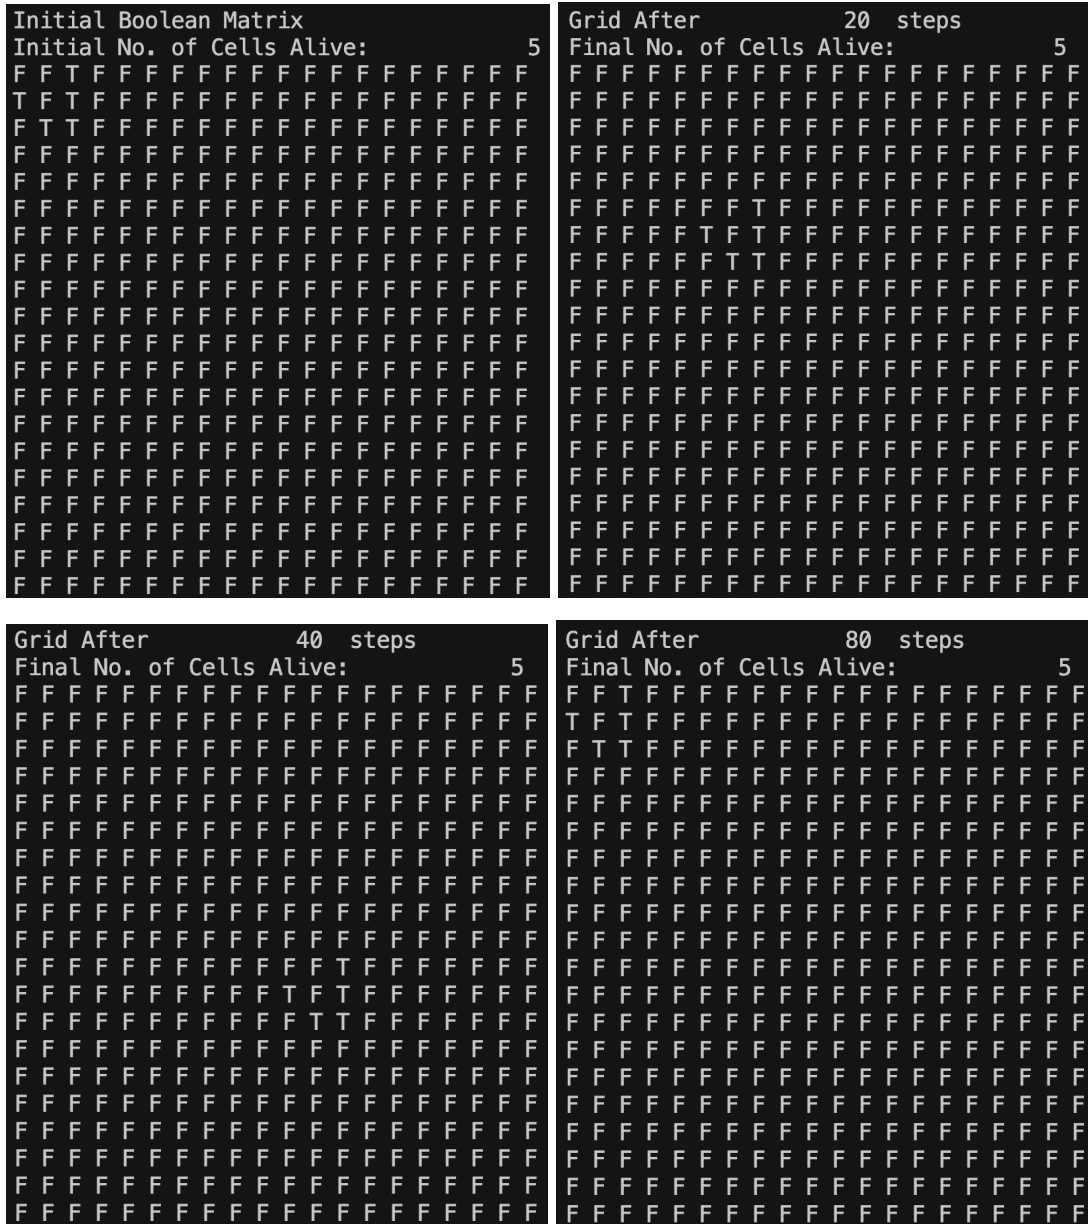
  call MPI_SENDRECV(send_up, local_ncols + 2, MPI_LOGICAL, up, 3, &
    recv_down, local_ncols + 2, MPI_LOGICAL, down, 3, &
    MPI_COMM_WORLD, MPI_STATUS_IGNORE, ierr)

  call MPI_SENDRECV(send_down, local_ncols + 2, MPI_LOGICAL, down, 4, &
    recv_up, local_ncols + 2, MPI_LOGICAL, up, 4, &
    MPI_COMM_WORLD, MPI_STATUS_IGNORE, ierr)

  local_game(:, 1) = recv_left
  local_game(:, local_ncols + 2) = recv_right
  local_game(1, :) = recv_up
  local_game(nrows + 2, :) = recv_down
end if

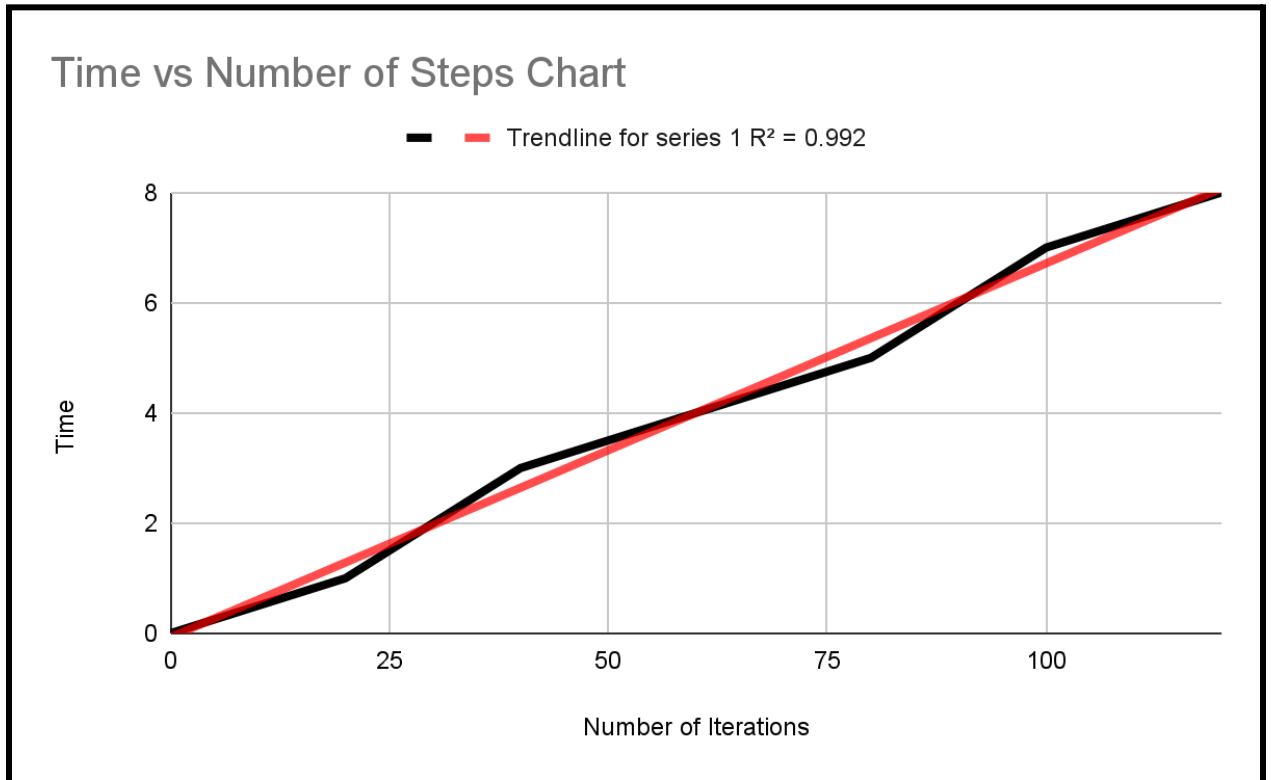
```

3. Here are my grids at steps 0, 20, 40, and 80 below:



I get the same pattern every 4 steps. The full pattern moves down and right 1 for every 4 steps. After each step is completed, I immediately adjust the periodic boundary conditions which allows for this pattern to be repeated. Since it is a 20 x 20, it takes 80 (20 x 4) steps to reach the exact same grid again. I also used 4 processors to do this.

4. Here's my plot:



This graph shows that there is a strong linear correlation between the execution time and the number of iterations. I used `system_clock` instead and iterations in increments of 20. The trendline equation is:

$Time = 0.0679 * No. of Iterations - 0.0714$ . The 0.0679 is very close to 0.05, so it is very close to a linear rate of an increase of 1 in time for every increase in 20 iterations. Startup cost here would technically be 0 and the cost per iteration would be 0.0679.